

Developing a Secure Password Manager with Customizable Cryptographic Primitives

Crypto Vlinh Team

February 22, 2026

Table of Contents

Abstract	5
1 Introduction	6
1.1 Context	6
1.2 Problem Statement	6
1.3 Objectives	6
2 Theoretical Background	8
2.1 Symmetric-key Encryption	8
2.1.1 The Advanced Encryption Standard (AES)	8
2.1.2 Cipher Block Chaining (CBC) Mode	8
2.1.3 Padding Mechanism	8
2.2 Key Derivation Functions	9
2.2.1 The Entropy Limitation of User Passwords	9
2.2.2 The PBKDF2 Algorithm	9
2.3 Authenticated Encryption	9
2.3.1 Data Integrity via HMAC	9
2.3.2 Analysis of Generic Composition Paradigms	10
2.3.3 Theoretical Superiority of Encrypt-then-MAC (EtM)	10
3 Methodology	11
3.1 Cryptographic Pipeline	11
3.1.1 System Flow Overview	11
3.1.2 Design Justification: Cryptographic Key Splitting	13
3.1.3 Encryption Phase	14
3.1.4 Authentication Phase	14
3.2 Vault Data Structure	14
3.2.1 Binary Layout Formatting	14
3.2.2 Safe Decryption and Verification Process	15
4 Results, Benchmarks, and Security Analysis	16
4.1 Implementation Performance and Benchmarking	16
4.1.1 Key Derivation Function (KDF) Latency	16
4.1.2 Systemic Efficiency	17
4.2 Security Analysis and Attack Resistance	17
4.2.1 Quantitative Security Metrics	17
4.2.2 Resistance to Brute-Force and Dictionary Attacks	18
4.2.3 Statistical Diffusion Analysis	18
4.2.4 Cryptographic Visualization via Hex Mapping	18

4.2.5 Integrity Protection and Resistance to Padding Oracle Attacks . .	20
5 Conclusion	21

List of Figures

3.1	Phase-based Encryption Pipeline	12
3.2	Phase-based Decryption Pipeline	13
3.3	Binary Vault File Format	15
4.1	KDF Latency vs. Iteration Count Performance Benchmark	16
4.2	Visual Evidence of Statistical Cryptographic Diffusion	18
4.3	Hexadecimal Mapping and Divergence Analysis	19

List of Tables

4.1	KDF Performance Benchmarking Data	17
4.2	Statistical Security Metrics for Cryptographic Pipeline	17

Abstract

The proliferation of digital services has significantly intensified the requirement for robust credential management. However, modern software development often relies on opaque, "black-box" cryptographic libraries that obscure underlying security mechanisms. This project addresses this challenge by designing and implementing a secure password manager that prioritizes transparency through the explicit orchestration of cryptographic primitives.

Utilizing a "white-box" implementation strategy, the system construction eschews high-level abstractions in favor of manual configuration of the Advanced Encryption Standard (AES) in Cipher Block Chaining (CBC) mode, combined with the Password-Based Key Derivation Function 2 (PBKDF2-HMAC-SHA256). To ensure both confidentiality and integrity, the architecture adopts the Encrypt-then-MAC (EtM) paradigm using HMAC-SHA256, effectively neutralizing integrity-based exploits such as padding oracle attacks.

Empirical evaluation demonstrates that the implementation achieves high security without compromising user experience. Security benchmarking indicates that a calibration of 390,000 PBKDF2 iterations results in a derivation latency of approximately 75.27 ms, providing a robust work factor against offline brute-force attempts. Furthermore, statistical analysis confirms near-ideal cryptographic properties, with a recorded avalanche effect of 50.10% and a Shannon entropy of 6.4133 bits per byte, signifying that the encrypted data is statistically indistinguishable from uniform noise. This project serves as a comprehensive validation of applying theoretical cryptographic standards to construct resilient, high-assurance security architectures.

Chapter 1

Introduction

1.1 Context

In the contemporary digital era, the proliferation of online services has unequivocally escalated the necessity for robust credential management. As individuals and organizations continually expand their digital footprints, the reliance on secure password managers has become a fundamental aspect of modern cybersecurity hygiene. Developing a comprehensive password manager serves as an archetypal paradigm for effectively bridging theoretical cryptographic concepts with practical information security applications.

1.2 Problem Statement

The pervasive risk of credential compromise stems largely from insecure storage mechanisms and flawed cryptographic implementations. A significant caveat in modern software development is the over-reliance on opaque, "black-box" cryptographic libraries. While these high-level abstractions offer convenience, they consistently obscure the underlying cryptographic mechanisms, thereby preventing developers from fully comprehending the nuances of data protection. To guarantee the robustness of a secure system, there is an imperative need to eschew these abstractions and instead implement and manage cryptographic primitives explicitly, facilitating exhaustive control and understanding of the encryption pipeline.

1.3 Objectives

This project aims to design and implement a secure password manager from the ground up, prioritizing transparency and rigorous adherence to cryptographic standards. The primary objective is the practical application of theoretical cryptography by manually orchestrating core cryptographic primitives, including the Advanced Encryption Standard (AES), Keyed-Hash Message Authentication Code (HMAC), and Password-Based Key Derivation Function 2 (PBKDF2), to construct a resilient local storage solution.

Furthermore, the project emphasizes a "white-box" implementation strategy. By explicitly constructing the algorithmic workflow, the system ensures absolute visibility and granular control over every phase of data processing. This transparent approach enables comprehensive security and performance evaluations. The system's cryptographic resilience is subsequently benchmarked and computationally analyzed against established

international standards, such as those promulgated by the National Institute of Standards and Technology (NIST) and Federal Information Processing Standards (FIPS). Ultimately, this endeavor aims to cultivate strict secure coding practices and foster a sophisticated methodology for designing high-assurance security architectures.

Chapter 2

Theoretical Background

This chapter establishes the fundamental cryptographic principles upon which the secure password manager is constructed. The theoretical framework presented herein is rigorously aligned with established international standards, specifically those promulgated by the National Institute of Standards and Technology (NIST) and the Internet Engineering Task Force (IETF).

2.1 Symmetric-key Encryption

2.1.1 The Advanced Encryption Standard (AES)

At the core of the system's confidentiality mechanism is the Advanced Encryption Standard (AES), standardized by NIST in FIPS PUB 197 [1]. AES, utilizing the Rijndael algorithm, operates on a Substitution-Permutation Network (SPN) structure. This architecture provides robust resistance against linear and differential cryptanalysis through iterated rounds of byte substitution, row shifting, column mixing, and round key addition, ensuring strong cryptographic diffusion and confusion.

2.1.2 Cipher Block Chaining (CBC) Mode

To encrypt data exceeding the standard 16-byte block size of AES, a mode of operation is required. The system implements Cipher Block Chaining (CBC) mode, as recommended in NIST SP 800-38A [2]. In CBC mode, each plaintext block is XORed with the preceding ciphertext block prior to encryption. This chaining mechanism ensures that identical plaintext blocks produce divergent ciphertext blocks. Crucially, the process necessitates an Initialization Vector (IV), which must be unpredictable (random) for each encryption operation. The use of a random IV in CBC mode guarantees "semantic security" against chosen-plaintext attacks, effectively concealing repetitive data patterns.

2.1.3 Padding Mechanism

As block ciphers like AES require input data length to be a strict multiple of the block size, a padding scheme must be employed. The Public-Key Cryptography Standards (PKCS) #7 scheme is utilized, ensuring that the plaintext consistently aligns with the 16-byte boundary prior to the AES-CBC encryption phase.

2.2 Key Derivation Functions

2.2.1 The Entropy Limitation of User Passwords

Passwords generated and memorized by human users inherently exhibit relatively low Shannon entropy. Consequently, they are fundamentally inadequate for direct application as cryptographic keys, which require a uniform distribution over the key space to resist brute-force attacks. To bridge this gap, NIST Special Publication 800-132 [3] mandates the utilization of a Password-Based Key Derivation Function (PBKDF) to cryptographically transform a low-entropy password into a high-entropy bit string suitable for symmetric encryption.

2.2.2 The PBKDF2 Algorithm

This system implements the PBKDF2 algorithm as specified in the Public-Key Cryptography Standards (PKCS) #5 v2.0 and subsequently published as RFC 2898 [4]. PBKDF2 can be formally defined as a function mapping a password (P), a salt (S), an iteration count (c), and a desired key length ($dkLen$) to a derived key (DK):

$$DK = PBKDF2(PRF, P, S, c, dkLen) \quad (2.1)$$

Where PRF denotes a pseudorandom function, typically a hash-based MAC. In this architecture, HMAC-SHA256 is selected as the optimal PRF . PBKDF2 introduces two pivotal cryptographic mechanisms to fortify the derived key:

- **Iteration Count for Key Stretching:** By recursively applying the PRF for c iterations (e.g., $c = 390,000$), PBKDF2 forces an artificial computational delay. Let $U_1 = PRF(P, S \parallel INT(i))$ and $U_c = PRF(P, U_{c-1})$. The block T_i is computed as the bitwise XOR of all U values: $T_i = U_1 \oplus U_2 \oplus \dots \oplus U_c$. This "key stretching" technique exponentially amplifies the computational work factor required for an adversary to test a single password guess, thereby severely mitigating the threat of brute-force and offline dictionary attacks without perceptibly impacting legitimate user authentication times.
- **Cryptographic Salt:** A cryptographically secure, uniformly random salt (S) of at least 128 bits is generated for each encryption instance. The salt acts as a unique non-secret parameter that fundamentally alters the input to the hash function. Mathematically, it guarantees that for any two identical passwords $P_1 = P_2$, if $S_1 \neq S_2$, then $DK_1 \neq DK_2$ with overwhelming probability. This mechanism completely neutralizes time-memory trade-off attacks, such as precomputed Rainbow Tables, by forcing the attacker to compute a distinct hash chain for every possible salt value.

2.3 Authenticated Encryption

2.3.1 Data Integrity via HMAC

While the AES algorithm in CBC mode guarantees data confidentiality, it provides negligible protection against active adversarial tampering. To ensure verifiable data integrity and authenticity, the system employs the Keyed-Hash Message Authentication

Code (HMAC) utilizing the SHA-256 hash function. Standardized in FIPS 198-1 [5], HMAC is formally defined as:

$$HMAC(K, m) = H((K \oplus opad) \parallel H((K \oplus ipad) \parallel m)) \quad (2.2)$$

Where H is the underlying cryptographic hash function (SHA-256), K is the secret authentication key derived securely via PBKDF2, m is the message, and $ipad$ and $opad$ are fixed padding constants. This nested construction mathematically precludes length-extension attacks that vulnerability traditional $H(K \parallel m)$ MAC designs.

2.3.2 Analysis of Generic Composition Paradigms

Establishing an intuitively secure channel necessitates the amalgamation of a symmetric encryption scheme ($\mathcal{SE}$) and a MAC scheme ($\mathcal{MA}$). The seminal formalized analysis by Bellare and Namprempre [6], subsequently corroborated by Krawczyk [7], dissects three generic composition paradigms:

1. *Encrypt-and-MAC (E&M)*: Compute $C = \mathcal{SE}_K(M)$ and $T = \mathcal{MA}_{K'}(M)$ independently. Transport (C, T) . Historically utilized in the SSH protocol.
2. *MAC-then-Encrypt (MtE)*: Compute a tag on the plaintext $T = \mathcal{MA}_{K'}(M)$, then encrypt the concatenation $C = \mathcal{SE}_K(M \parallel T)$. Transport C . This paradigm was the foundation of SSL and early TLS specifications.
3. *Encrypt-then-MAC (EtM)*: Encrypt the plaintext $C = \mathcal{SE}_K(M)$, then compute the tag over the resulting ciphertext $T = \mathcal{MA}_{K'}(C)$. Transport (C, T) . Adopted by IPsec.

2.3.3 Theoretical Superiority of Encrypt-then-MAC (EtM)

Rigorous cryptographic proofs establish that **Encrypt-then-MAC (EtM)** is the singular composition paradigm guaranteed to provide chosen-ciphertext attack (IND-CCA) security and Integrity of Ciphertexts (INT-CTXT), provided the underlying encryption is IND-CPA secure and the MAC is strongly unforgeable (SUF-CMA) [6].

Conversely, the MtE paradigm—despite its intuitive appeal of authenticating only valid plaintexts—was definitively proven to be generically insecure by Krawczyk [7]. When applied to block ciphers operating in CBC mode, MtE vulnerabilities invariably lead to devastating practical exploits, most notably the "padding oracle" attack (e.g., the Vaudenay attack).

By implementing the EtM architecture, this password manager inherently circumvents these systemic vulnerabilities. The cryptographic pipeline mandates that the HMAC tag is fully validated against the ciphertext and Initialization Vector *before* any decryption routine is initiated. If the tag validation fails, the ciphertext is instantly discarded without processing, thereby mathematically denying the adversary any oracle access to the underlying decryption or padding mechanisms.

Chapter 3

Methodology

This chapter delineates the practical implementation of the theoretical cryptographic primitives discussed in Chapter 2. The methodology section details the concrete architecture of the password manager, explicitly defining how AES, HMAC, and PBKDF2 are orchestrated to construct a secure computational pipeline and a robust data storage format.

3.1 Cryptographic Pipeline

3.1.1 System Flow Overview

The core operational logic of the password manager is deterministically driven by the user's Master Password. Figure 3.1 illustrates the encryption pipeline, transitioning from the initial password to the final secure vault. The process initiates with PBKDF2 deriving a Master Key, which is subsequently split into independent keys for encryption and authentication.




Figure 3.1: Phase-based Encryption Pipeline

Conversely, Figure 3.2 demonstrates the decryption pipeline. This routine rigidly

enforces the security invariants established during encryption by verifying the authentication tag prior to any decryption attempt.




Figure 3.2: Phase-based Decryption Pipeline

3.1.2 Design Justification: Cryptographic Key Splitting

Following the derivation of a high-entropy Master Key (MK) via PBKDF2, a critical architectural decision is executed: **Key Splitting**. The 64-byte MK is bisected to produce two mathematically independent 32-byte keys: an AES Key (K_{enc}) for confidentiality and an HMAC Key (K_{mac}) for integrity.

This separation of concerns is a fundamental requirement for secure system design.

As formally analyzed by Bellare and Namprempre [6], deploying the *same* cryptographic key for both the symmetric encryption scheme and the Message Authentication Code (MAC) inherently invalidates the generic security proofs of the composition. Reusing a key introduces perilous cross-primitive correlations, potentially allowing an adversary to leverage weaknesses or state from the MAC function to compromise the block cipher, or vice versa. By deriving K_{enc} and K_{mac} independently, the system guarantees that the confidentiality and integrity mechanisms remain cryptographically isolated and secure.

3.1.3 Encryption Phase

The encryption phase transforms the serialized vault data from plaintext into an opaque ciphertext through a series of deterministic transformations. Initially, the system utilizes a cryptographically secure pseudo-random number generator (CSPRNG) to produce a unique 16-byte Initialization Vector (IV), ensuring that identical plaintexts result in distinct ciphertexts over multiple encryption cycles. Before the cryptographic core is engaged, the plaintext payload is processed using the PKCS#7 specification, adding defensive padding to align the data length with the 16-byte block size mandated by the cipher. Finally, the AES algorithm is executed in Cipher Block Chaining (CBC) mode, utilizing the dedicated encryption key and the generated IV to yield the final ciphertext (C).

3.1.4 Authentication Phase

To achieve comprehensive security against chosen-ciphertext attacks, the system strictly implements the Encrypt-then-MAC (EtM) paradigm. Following the ciphertext production, a Message Authentication Code is generated to bind the encrypted data to its specific context. This is accomplished by concatenating the IV with the ciphertext (C) to form a comprehensive authentication payload. The HMAC-SHA256 algorithm is then applied to this payload using the dedicated K_{mac} , producing a 32-byte authentication tag (Tag). This architecture ensures the integrity of the entire encrypted package, allowing the system to detect any unauthorized modifications to either the ciphertext or the IV before the decryption logic is ever exercised.

3.2 Vault Data Structure

3.2.1 Binary Layout Formatting

The output of the cryptographic pipeline is serialized into a rigid binary format to ensure interoperability and parameter persistence as visualized in Figure 3.3. This structure adheres to the security recommendations of NIST SP 800-132 regarding the adjacent storage of KDF parameters and protected data.




Figure 3.3: Binary Vault File Format

The vault file consists of four contiguous segments: a 16-byte random Salt for key derivation, a 16-byte IV for the block cipher, and a 32-byte HMAC-SHA256 Tag for integrity verification. These parameters precede the variable-length ciphertext, facilitating a linear parsing routine that can reconstruct the cryptographic state required for safe data retrieval.

3.2.2 Safe Decryption and Verification Process

The retrieval of plaintext data from the vault hinges upon the uncompromising principles of the EtM composition, which dictates that authentication must precede any decryption transformation. The process begins with the reconstruction of the Master Key (MK) using the retrieved Salt and the user's password, followed by a secondary key splitting to recover K_{enc} and K_{mac} . Before the AES cipher is initialized, the system computes a candidate MAC from the stored IV and ciphertext.

A constant-time comparison is then performed between the candidate and the stored Tag . In the event of a verification failure, the system immediately rejects the vault as compromised, aborting the routine without ever invoking the AES decryption function. This critical invariant completely neutralizes potential vulnerabilities to padding oracle attacks. Only upon a successful integrity match is the AES-128-CBC decryption initialized, transforming the ciphertext back into plaintext and stripping the PKCS#7 padding to recover the original data.

Chapter 4

Results, Benchmarks, and Security Analysis

This chapter presents a comprehensive evaluation of the secure password manager implementation, focusing on both performance benchmarks and rigorous security analysis. The empirical results demonstrate the effectiveness of the selected cryptographic primitives and the robustness of the system against common architectural vulnerabilities and statistical cryptanalysis.

4.1 Implementation Performance and Benchmarking

4.1.1 Key Derivation Function (KDF) Latency

The performance of the system is primarily characterized by the computational overhead of the PBKDF2-HMAC-SHA256 key derivation process. Following the security guidelines defined in NIST Special Publication 800-132 [3], the iteration count was calibrated to maximize the work factor for an adversary while maintaining an acceptable user experience. Experimental results, visualized in Figure 4.1, indicate that an iteration count of 390,000 provides a derivation latency significantly below the one-second threshold recommended for interactive applications.




Figure 4.1: KDF Latency vs. Iteration Count Performance Benchmark

The empirical timing data for various iteration counts is summarized in Table 4.1. The results indicate that even at the maximum tested value of 1,000,000 iterations, the latency remains well under 200 milliseconds, demonstrating that the system provides high security with negligible impact on user-perceived performance.

Table 4.1: KDF Performance Benchmarking Data

Iteration Count	Time (ms)
1,000	0.23
10,000	2.02
100,000	18.69
400,000 (Recommended)	75.27
1,000,000	186.25

This artificial delay serves as a pivotal defense mechanism, exponentially increasing the cost of offline brute-force attempts by requiring a substantial investment in computational time for every password guess tested. As shown in the benchmarking plot and table, the relationship between iteration count and time is linear, allowing for predictable security scaling based on the targeted hardware profiles.

4.1.2 Systemic Efficiency

Beyond the initial key derivation phase, the operational efficiency of the AES-128-CBC encryption and HMAC-SHA256 authentication routines remains exceptionally high. The overhead introduced by the Encrypt-then-MAC (EtM) composition is negligible compared to the cryptographic strength it provides. Data throughput for vault operations is limited only by the underlying file system I/O, ensuring that the security measures do not impose a perceptible performance penalty during standard usage cycles.

4.2 Security Analysis and Attack Resistance

4.2.1 Quantitative Security Metrics

To validate the cryptographic integrity of the system, a series of statistical analyses were conducted over 100 random trials. These tests focused on the randomness and diffusion properties of the generated ciphertext. The primary metrics, including Shannon entropy and the statistical avalanche effect, are summarized in Table 4.2.

Table 4.2: Statistical Security Metrics for Cryptographic Pipeline

Metric	Value	Standard Deviation (σ)
Average Shannon Entropy	6.4133 bits/byte	± 0.1184
Statistical Avalanche Effect	50.10%	$\pm 1.52\%$
Number of Trials	100	N/A

The empirical findings confirm that the implementation adheres closely to the theoretical ideals of modern cryptography, providing a robust foundation for secure data storage.

4.2.2 Resistance to Brute-Force and Dictionary Attacks

The architectural defense against unauthorized access is rooted in the rigorous application of PBKDF2 with a unique, cryptographically secure 128-bit salt for every vault instance. As mandated by RFC 2898 [4], the use of random salts completely neutralizes the threat of precomputed Rainbow Table attacks by ensuring that identical master passwords result in unique derived keys. Furthermore, the high iteration count effectively mitigates dictionary and brute-force attacks by drastically reducing the rate at which an adversary can execute trial authentications, thereby shifting the security margin in favor of the defender.

4.2.3 Statistical Diffusion Analysis

The quality of the encryption is validated through the **Avalanche Effect**, representing the principle—stipulated in the AES specification (FIPS 197) [1]—that a single-bit change in the plaintext should result in an average of 50% change in the ciphertext bits. The recorded average of 50.10% signifies near-ideal cryptographic diffusion. This property ensures that no discernible patterns from the original plaintext are preserved in the encrypted output, effectively neutralizing frequency analysis attacks. Visual proof of this diffusion property is presented in Figure 4.2.




Figure 4.2: Visual Evidence of Statistical Cryptographic Diffusion

The **Shannon Entropy** of 6.4133 bits per byte further corroborates these findings. This value indicates that the ciphertext is statistically indistinguishable from uniform noise, ensuring that the protected data remains opaque to any non-adversarial observation or pattern-matching algorithm.

4.2.4 Cryptographic Visualization via Hex Mapping

The spatial distribution of data within the vault is visualized through annotated hex maps. Figure 4.3 provides a detailed side-by-side comparison of the binary structure and the non-deterministic nature of the cryptographic pipeline.



(a) Annotated layout of a single vault

(b) Comparison demonstrating semantic security

Figure 4.3: Hexadecimal Mapping and Divergence Analysis

Figure 4.3a identifies the segments dedicated to the Salt, IV, and HMAC Tag. The non-deterministic nature of the system is further demonstrated by comparing two vaults generated with the same master password. Despite identical inputs, the cryptographic pipeline produces completely divergent binary outputs due to the unique Salt and IV parameters. This divergence, as depicted in Figure 4.3b, serves as empirical proof of the system’s ability to maintain semantic security.

4.2.5 Integrity Protection and Resistance to Padding Oracle Attacks

The implementation's most critical security property is its adherence to the Encrypt-then-MAC (EtM) paradigm. By authenticating the ciphertext and IV prior to decryption, the system achieves Integrity of Ciphertexts (INT-CTXT). This architectural design, formally analyzed by Bellare and Namprempre [6], effectively neutralizes many classes of active attacks, including the devastating Padding Oracle exploit described by Hugo Krawczyk [7]. Any unauthorized modification to the vault file is detected during the HMAC validation phase, causing the system to terminate the process before any decryption or padding removal is attempted.

Chapter 5

Conclusion

The design and implementation of this secure password manager have successfully demonstrated the practical application of core cryptographic primitives within a cohesive security architecture. By transitioning from abstract theoretical concepts to a concrete “white-box” implementation, the project has achieved its primary objective of constructing a transparent and auditable encryption pipeline. The integration of AES-128 in CBC mode, paired with HMAC-SHA256 in an Encrypt-then-MAC configuration, provides a high-assurance solution that guarantees both the confidentiality and integrity of stored credentials while resisting sophisticated chosen-ciphertext attacks.

Beyond its technical functionality, the system serves as an empirical validation of international security standards, specifically those promulgated by NIST and the IETF. The rigorous benchmarking of the PBKDF2 algorithm ensured an optimal balance between computational work factor and user experience, while statistical analysis of the vault structure confirmed the existence of robust cryptographic diffusion. This project underscores the critical importance of defensive design patterns, such as key splitting and the strict separation of encryption and authentication layers, in building resilient information systems.

Beyond its development, the system’s robust security posture provides a foundation for several avenues of future enhancement. A primary area for improvement lies in the transition of the key derivation mechanism to more modern, memory-hard functions. Specifically, replacing PBKDF2 with Argon2id—the winner of the Password Hashing Competition—would provide significantly greater resistance against hardware-accelerated brute-force attacks from GPUs and ASICs. Unlike PBKDF2, which is primarily CPU-bound, Argon2id mandates a configurable amount of memory, effectively leveling the playing field between an attacker and a legitimate user.

Furthermore, future iterations of the software could implement advanced memory protection techniques to safeguard sensitive data during runtime. While the current architecture ensures “security at rest,” the implementation of memory-zeroing routines and the encryption of in-RAM data structures would provide “security in use,” mitigating the risks associated with cold-boot attacks or unauthorized memory dumps. Additionally, the integration of multi-factor authentication (MFA) and the transition toward a decentralized synchronization model could further elevate the system’s utility and security profile in a distributed environment.

References

- [1] National Institute of Standards and Technology, “Advanced encryption standard (aes),” Tech. Rep. FIPS PUB 197, U.S. Department of Commerce, November 2001.
- [2] M. Dworkin, “Recommendation for block cipher modes of operation: Methods and techniques,” Tech. Rep. SP 800-38A, National Institute of Standards and Technology, December 2001.
- [3] M. S. Turan, E. Barker, W. Burr, and L. Chen, “Recommendation for password-based key derivation: Part 1: Storage applications,” Tech. Rep. SP 800-132, National Institute of Standards and Technology, December 2010.
- [4] B. Kaliski, “Pkcs #5: Password-based cryptography specification version 2.0.” RFC 2898, 2000.
- [5] National Institute of Standards and Technology, “The keyed-hash message authentication code (hmac),” Tech. Rep. FIPS PUB 198-1, U.S. Department of Commerce, July 2008.
- [6] M. Bellare and C. Namprempre, “Authenticated encryption: Relations among notions and analysis of the generic composition paradigm,” in *Advances in Cryptology — ASIACRYPT 2000*, pp. 531–545, Springer, 2000.
- [7] H. Krawczyk, “The order of encryption and authentication for protecting communications (or: How secure is ssl?),” in *Advances in Cryptology — CRYPTO 2001*, pp. 310–331, Springer, 2001.